Assignment 4: Inverse Kinematics and Trajectories

Robot Kinematics and Dynamics

Prof. Howie Choset

Contents

1   Overview

This assignment reinforces the following topics:
- Analytic IK
- Numerical IK
- Trajectories for Robot Control

2   Background

2.1   Analytic IK

One approach to IK is to solve for joint angles given a workspace end effector position manually, via algebraic methods or geometric analysis, for each different type of manipulator. This is very different from FK, which is pretty much formulaic, going through the transforms of each link to get the end effector position. The range of arms for which analytical solutions exist are only arms that are fully constrained, and solutions generally have to be further manually inspected to check for an arm exceeding joint limits or other practical considerations.

Analytical IK has the advantage that it is precise, fast (once you figure out the solution), and global, meaning that you can always find a solution for joint angles given an end effector position that is within the arm's workspace. So in short, with analytical IK:

- If there is a solution(s), you will find it.
- If there is no solution, you will know that.

A more flexible method we will discuss below is numerical IK, which can work around most of the limitations posed by things like redundancy and joint limits. However, it is approximate and local, meaning that you are not always guaranteed to find a 'good' solution.

2.2   Numerical IK

An alternative approach to analytical IK is numerical IK, or optimization-based IK. In this case we use the forward kinematics and a cost function to try to find a solution to the inverse kinematics. The general idea is to design a function that takes in a desired end effector pose and produces a scalar 'cost' (often an error based on distance to the desired pose) that an optimizer will try to minimize by varying the input variables. In this case the input variables are joint angles which are used to calculate the forward kinematics as part of the cost function.

Overall the process works like:

1. Pick a desired end effector pose

2. Pick an initial arm configuration (this is important detail; usually you are already in a configuration that you wish to start from)

3. Evaluate the forward kinematics at the current configuration

4. Evaluate the update rule for the input variables (e.g., gradient descent, a method based on the gradient of the cost function)

5. Check for 'stopping conditions', often based on the change in cost or the number of iterations (e.g., for gradient descent this can be related to the magnitude of the gradient); if we meet these stopping conditions, then return the current value of the input variables as our solution

6. Update the input variables and go back to step 3

Different optimizers have different methods of picking the next configuration to try, but most are based on gradient methods. Essentially, they try to figure out the slope of the cost function and trying to go 'downhill' until they reach a minimum.

There are distinct advantages and disadvantages of numerical IK compared to analytical IK.

2.2.1 Advantages

- Numerical IK is a 'black box'. As long as you have forward kinematics for an arm, you can make a cost function, pick an optimizer, and solve the inverse kinematics. It's general and relatively straightforward to implement.
- Numerical IK can handle arms that are under-constrained (redundant) or over-constrained.
- You can modify the cost function for numerical IK to include whatever you want. Joint limits, knowledge of obstacles, other high-level requirements, etc. can all be added relatively easily.

2.2.2 Disadvantages

- Numerical IK is a 'black box'. You do not necessarily get any intuition into the underlying structure of the problem. There might be multiple solutions to a desired pose, or none at all. You have to be careful with how you interpret and use the result that an optimizer spits out.
- Solutions are local, and subject to initial conditions. Optimizers all need an initial guess on where to start, and this guess can dramatically affect the solution the optimizer finds. The most reliable way to use numerical IK for arms is to start in a configuration that is as close as possible to the desired configuration, essentially 'warm-starting' the optimizer instead of searching over large areas of the configuration space.
- Numerical IK is usually less efficient computationally than analytical IK. Frequently this is not the limiting factor for controlling a robot. However, there are some cases where you might want to sample the workspace a huge number of times (one example of this in robotics research is the *RRT*, or Rapidly exploring Random Tree) and in such cases the speed of your IK might be the limiting factor.

2.3 Trajectories

When controlling a robotic arm, it is usually important to handle things in a way that is more sophisticated than just 'banging' the arm from one configuration to another. The general idea is to not command anything a robot physically do. Although we have not yet discussed dynamics in the class, dynamics affect the ability to control the position or velocity of any manipulator in the physical world. In short:
- Real robots have mass, m
- $F = ma$
- F is limited, so we should be careful about our desired a.

One way to mitigate effects of the dynamic limits of an arm is to use *trajectories* to smoothly move the arm from one configuration to another. (When thinking about trajectories, you can pose

them as trajectories in work space or configurations space. Here we'll just think of trajectories in configuration space, but in the code problems you'll work with workspace trajectories as well.)

### 2.3.1 Bang-Bang Control (No Trajectories)

The simplest method is to ignore the pesky bounds of reality and rely on the underlying controller of the robot's actuators to 'catch up' to your commands. Depending on the specific application, and they way the low-level controllers are tuned, this can actually work. However, it has the disadvantage that the robot's motion is less predictable since the low-level controllers dictate the details of how the robot actually moves between the start and goal position – which is something that you often want to have control over! Additionally, with this approach you are usually 'jerking' the arm aggressively which can be dangerous and put unnecessary wear and tear on the arm's joints.

2.3.2 Constant Velocity (Better Than Nothing)

The next-simplest method you might think of would be to respect the reality that a robot's joints can't move with infinite speed. Looking at the case of moving a single DOF between two angles, we can choose an amount of time to take to move between the angles and command a corresponding constant velocity.

$T_{\text{move}} = (\theta_{\text{end}} - \theta_{\text{start}})/V_{\text{cruise}}$

This is a step in the right direction, but it has issues where we start and stop. To see the problem, let's think about the position, velocity, and acceleration.
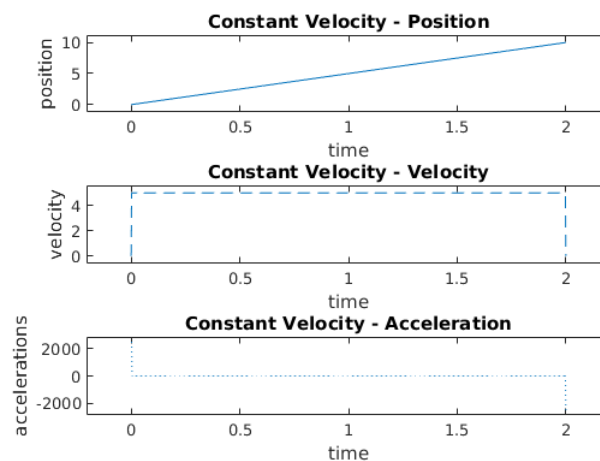

Because we jump from stop to a constant velocity and back again, at the ending steps we see big spikes in acceleration. In theory, these spikes are infinite acceleration, and in practice we will see some large finite value depending on how finely we discretize time (size of our Δt in the simulation or arm controller). Large accelerations are undesirable because if the arm has any mass, then we are commanding large torques that are usually outside the capabilities of the robot. Again, the goal of using trajectories is to always command things the arm can physically do, so that it is playing 'catch up' as little as possible.

3 Written Questions

For all of the following questions, please fully compute all answers and show your work for full credit, unless otherwise specified. Submit your answers in a PDF entitled `writeup.pdf` in your handin directory. All answers must be typed, but diagrams may be hand-drawn and scanned in. However, they must be tidy and fully legible! Consider drawing them in a black or blue pen. All units are in radians and meters, where appropriate.

1) Analytic IK: RPR robot
   Consider the following robot
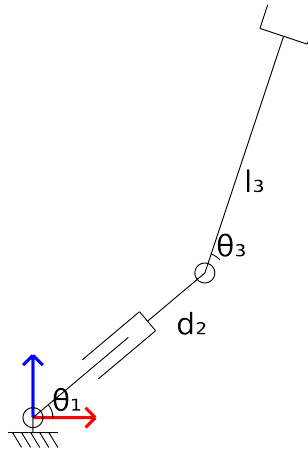
(Define the entire distance between the revolute joints as d_2).

(1) [7 points] Geometrically solve for θ_1, d_2, and θ_3 which result in an end-effector pose of $\begin{bmatrix} x_e \\ y_e \\ \theta_e \end{bmatrix}$. Show your work! HINT: Break this problem down into an easier subproblem; are there known workspace positions of other parts of the robot, given the end effector pose?

(2) [3 points] Assume there are no joint limits, except that $d_2 >= 0$. What (if any) limitations are there on achievable end-effector poses $\begin{bmatrix} x_e \\ y_e \\ \theta_e \end{bmatrix}$. Explain your answer in 1-2 sentences.

2) Numerical IK: Cost function
   Consider the following robot
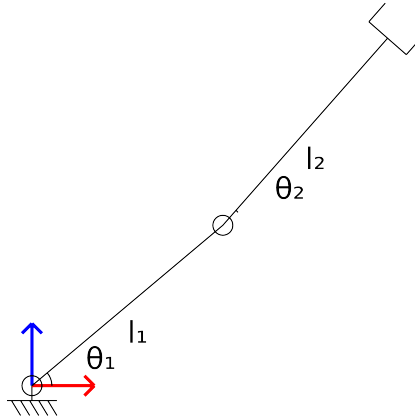
(1) [3 points] The goal is to find $\theta_d = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix}$ which achieves an end effector position $p_e = \begin{bmatrix} f_x(\theta) \\ f_y(\theta) \end{bmatrix}$ (for forward kinematics f and joint configuration θ) as close as possible to the desired position $p_d = \begin{bmatrix} x_d \\ y_d \end{bmatrix}$. Let cost function h be the square of the distance between p_e and p_d. Write this cost function in terms of the forward kinematics function $f(\theta)$ and p_e (e.g. you may use f_x and f_y, but do not use l_1 or l_2).

(2) [2 points] Write the gradient of h, $\nabla h(\theta) = [\frac{\partial h}{\partial \theta_1}, \frac{\partial h}{\partial \theta_2}]^T$, in terms of f, p, and partial derivatives of f with respect to θ_i. Note: do not actually write out the derivatives of f; just use, for example, $\frac{\partial f}{\partial \theta_1}$. HINT: if you are not familiar with vector differentiation, first write out $h(\theta)$ as a scalar expression by multiplying the individual vectors, and then compute the required partial derivatives as given above in the definition of $\nabla h(\theta)$.

(3) [3 points] Reorder these terms to write this expression in terms of the end effector Jacobian, J, along with f and p.

(4) [2 points] What would the expression for the gradient from the previous part look like for a 100 link serial chain of revolute modules? Is there any difference? Describe why or why not.

4   Feedback

1) Feedback Form   *5 points*

   We are always looking to improve the class! To that end, we're looking for your feedback on the assignments. When you've completed the assignment, please fill out the feedback form.

5   Code Questions

Copy the Code Handout folder to some location of your choice. Open Matlab and navigate to that location. Whenever you work on the assignment, go into this directory and run `setup.m`.

We will extend the Robot class you created in the last assignment; begin by copying this `Robot.m` file into the root of the Code Handout folder.

1) Numerical Jacobian

(1) [15 points] **Implementation** The first task in this assignment is to write the last Jacobian you'll ever need – a Jacobian that can work for any frame generated by the forward kinematics. You will fill in a `jacobians` function in the `Robot.m` file; copy and paste from the template in the `jacobian_source.txt` file.

This function computes a 3xn Jacobian (where n is the number of joints) for each of the $n + 1$ frames in the diagram.

The code describes the conventions used for the arguments and result.

Whereas in previous assignments you have calculated the analytic derivative for each element of the Jacobian, here you will use a numerical approach.

Given a general function $g(x)$, one can compute a numerical approximation of the derivative $g'(x)$ by computing $(g(x + dx) - g(x - dx))/(2 * dx)$.

In the code, we will take the same approach for each element of the Jacobian. (Note – as with the analytical approach, the rotation fields are computed by inspection for now). For example, for the element of the Jacobian corresponding to the x value for frame i and joint j (given a joint configuration $\mathbf{Q}$), one can evaluate the forward kinematics f_x at $\mathbf{Q} + dq_j$ and at $\mathbf{Q} - dq_j$, and divide the difference by 2ϵ (where dq_j is a vector of zeros with ϵ for element j).

(2) [0 points] **Verification** To verify the results, run the `sample_velocities` function, and verify that (1) on the first plot, the paths match *and* the end effector velocity vectors are tangent to the path that it takes, and (2) on the second plot, that the x and y velocities match.

Also, use this function to find the Jacobian for an RRR arm at configuration $[0, 0, 0]^T$. Look at the third row of the Jacobian for each frame – the effect of joint angle velocity on the orientation of the frame. For the first frame (recall, attached to the first link), only the first joint angle should have an effect (e.g., this row should be $[1, 0, 0]$). For the second frame, the first and second joint angles should have an effect. The third and fourth frames (third link and end effector, respectively) should each be affected by the motion of every joint.

Note: you do not have to submit anything for this part; however, you should use this to detect any problems with your code before continuing!

(3) [5 points] **Application** Because we have a Jacobian for any planar revolute-joint based arm, we can now address problems that would have been much more difficult before. Imagine a RRRRRRRRRR arm (10 R's), with links each equal to 5. Using the robot

class you have written, compute the joint torques necessary to apply a wrench[1] of $\begin{bmatrix} 0 \\ 5 \\ 2 \end{bmatrix}$
with the end effector when the robot is at the following configuration:

$[0, 0.1, 0.2, 0.3, 0.4, 0, -0.1, -0.2, -0.3, -0.4]^T$

Write the MATLAB expressions to compute this result (e.g., creating a robot object, calling the jacobians function, and multiplying the result) in a file titled `jacobian_example.m` in the `ex_01` directory.

(Just for fun, imagine the number of terms for the analytic Jacobian for this arm. If you could write that out the very first time without making a mistake or a sign error, you might be the first person to ever do so...)

2) Numerical IK

(1) [15 points] **Implementation** The next extension to the Robot class will be a function to compute inverse kinematics using gradient descent, as discussed in lecture. You will fill in an `inverse_kinematics` function in the Robot.m file; copy and paste from the template in the `inverse_kinematics_template.txt` file.

This function accepts a vector of joint angles for initial conditions, and a desired end effector location. It returns the resulting joint angles that are the solution of the gradient descent process.

The code describes the conventions used for the arguments and result.

For this problem, assume the cost function is the sum squared error between the end effector and the goal point, ignoring joint angle limits. Use the solution to the second written problem to help compute the gradient of the cost function, and use a loop to continue updating the joint angles until no significant progress is made. REMINDER: You use the *negative* gradient in this process, multiplied by a small stepsize α.

Your code should handle two cases:

- If the goal is a 2x1 vector, it should only consider $[x, y]^T$ position of the end effector.
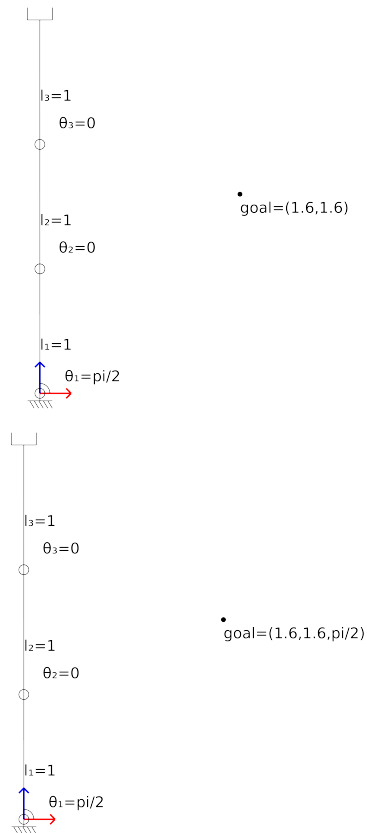- If the goal is a 3x1 vector, it should consider $[x, y, \theta]^T$ pose of the end effector.

HINT 1: Note that the `end_effector` function in the robot class returns the end effector's $[x, y, \theta]$ pose. This will probably be very useful in coding the equation!

HINT 2: For the 3x1 goal vector case, note that the form of the solution does not change; you must simply use an additional row of the Jacobian and include the end effector's orientation.

(2) [0 points] **Verification** To verify the results, solve the examples in the following figures:

---
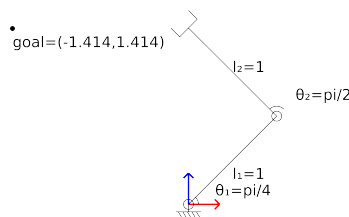
[1]Recall this is a combined force and torque vector; the first and second terms are forces in x and y, and the third term is a moment.

$l_3=1$
$\theta_3=0$

goal=(1.6,1.6)

$l_2=1$
$\theta_2=0$

$l_1=1$
$\theta_1=pi/2$

$l_3=1$
$\theta_3=0$

goal=(1.6,1.6,pi/2)

$l_2=1$
$\theta_2=0$

$l_1=1$
$\theta_1=pi/2$

Note: you do not have to submit anything for this part; however, you should use this detect any problems in your code before continuing!

Use the listed initial configurations, and compute the joint angles. For example, for the first problem, this would be:

```
robot = Robot([1;1],[1;1],[1;1],1);
initial_angles = [pi/4; pi/2];
goal = [-sqrt(2); sqrt(2)];
final_angles = robot.inverse_kinematics(initial_angles, goal);
```

For these three problems, the resulting joint angles should achieve the goal end effector position when passed into the forward kinematics function; you can test by running

```
robot.end_effector(final_angles)
```

Note that for examples 2 and 3, the resulting joint angles are different because of the inclusion of the end effector orientation constraint.

Again, a caution here is that numerical methods for IK are *local* - for example, the first problem above has two solutions (elbow up/elbow down), but numerical IK only finds the closest one to the initial configuration (or in the case of constraints such as joint angles, may not find any valid solution). This is important to consider for real-world applications.

This local approach can be a benefit over analytical methods, though: for example, in the first problem, if the goal was $[-2, 2]^T$, analytic IK would return no solution, whereas numerical methods would still minimize the cost as best as possible, and result in a "close" solution.

Another note: you can make the optimization more accurate by decreasing the step size and stopping condition; however, for our 'plain' gradient descent implementation this can result in long runtimes. There are built-in numerical optimization methods in MATLAB that are much more efficient than our simple approach here, resulting in more accurate results in less time.

3) Trajectories

When controlling a robot, one option is to just command a destination joint angle configuration, and allow the robot's internal joint angle controllers to move the joints towards this goal. However, depending on the starting position of the robot, this can result in large dangerous motions, and also will not give you any control over *how* the robot gets to this destination. Are there obstacles you want the robot to avoid? Do you want the resulting motion to be smooth?

Instead, one can define a *joint angle trajectory* for the robot's joints to follow through this motion to the goal point, and then iteratively command points along this trajectory to result in a smooth, controlled motion along a well defined path.

(1) [5 points] **Straight Line Joint Tractories**

Given a starting joint configuration and an goal joint configuration one option is to connect the start and ending joint angles with a straight line in *joint space*. This means that you linearly interpolate between the start and end joint angles[2].

In this problem, you will fill in `linear_joint_trajectory.m`. This function takes a column vector of start angles, a column vector of end angles, and a trajectory resolution (number of points). It linearly interpolates between the angles in order to reach the end points at the given time. It will return a matrix of n angles by t timesteps.

(2) [0 points] **Verification**

For an RRR arm with unit link lengths, interpolate between the start state $\begin{bmatrix} \pi/4 \\ \pi/2 \\ \pi/2 \end{bmatrix}$ and the end state $\begin{bmatrix} 0 \\ \pi/2 \\ 0 \end{bmatrix}$, with 100 steps.

Run the `visualize_trajectory.m` function with the relevant robot and resulting trajectory matrix as the argument. The joint angle plots should be straight lines, as this is a linear interpolation in joint space. What does the end effector path look like? Does this make sense?

Note: you do not have to submit anything for this part; however, you should use this detect any problems in your code before continuing!

(3) [10 points] **Linear Workspace Joint Trajectories**

---

[2]In MATLAB, the function *linspace* will linearly interpolate between two numbers with a given resolution

Now assume that you have a start joint configuration of the robot, and want to reach a workspace goal position g by moving in a straight line in *the workspace*. This can be useful for moving between obstacles, or to approach obstacles for manipulation.

In this problem, you will fill in `linear_workspace_trajectory.m`, which takes a robot object, a column vector of start angles, a workspace goal position (it should accept points in both $\mathbb{R}^2$ and SE(2)), and a trajectory resolution (number of points). This function should return a matrix of n angles by t timesteps which result in a straight line workspace trajectory between the inputs.

The general approach to this problem is to first linearly interpolate between initial and goal workspace positions, and then solve the IK for each of these positions in turn. Note that when solving for the i^{th} column of the result trajectory matrix, you should use the $(i-1)^{\text{th}}$ column as the initial condition for the numerical IK, as this will ensure the resulting trajectory is smooth.

(4) [0 points] **Verification**

We will repeat the verification from the last part of this problem, except we use the workspace position for the end point instead of the joint angles:

For an RRR arm with unit link lengths, interpolate between the start state $\begin{bmatrix} \pi/4 \\ \pi/2 \\ \pi/2 \end{bmatrix}$ and the end workspace position $\begin{bmatrix} 1 \\ 2 \end{bmatrix}$, with 100 steps.

Run the `visualize_trajectory.m` function with the relevant robot and resulting trajectory matrix as the argument. Is the end effector workspace path a straight line? What about the joint angle plots? Does this make sense?

Note: you do not have to submit anything for this part; however, you should use this detect any problems in your code before continuing!

Additional Notes: End goal   You may have noticed that the `linear_joint_trajectory.m` file had a second argument of a vector of joint angles, and the `linear_workspace_trajectory.m` file has a second argument of a a workspace position.

Although this could be done in different ways, often the first function would be used when moving between joint "waypoints," rather than between a current robot position and a desired workspace position. You will see an example of this in the "teach and repeat" problem in the hands-on part of this assignment.

The workspace trajectory would be difficult to complete if it were between two joint configurations, because one would have to enforce a continuous constraint on the joint angles in addition to the linear workspace constraint; this is not in general possible (as a concrete example, imagine an RR arm moving from an "elbow up" to an "elbow down" configuration where the end effector could not move (e.g., a straight line between two identical points in the workspace).

6 Hands-On Questions

For this lab, there are two different robots (one with 2DOF and one with 3DOF). You will be expected to demonstrate your code to the TAs, and you should be prepared to demonstrate on either robot.

You will use your code from the previous section for this lab, and therefore will need to create 'Robot' objects to pass into the functions below. The dimensions of the links will be given on the table beside the robots; the link/joint/end effector masses are not used for this lab and can be ignored (i.e., just set to '0' or '1').

1) "Teach/Repeat"

One simple way to control the behavior of real robots is to skip the pesky IK step, and just demonstrate the joint angles that you wish the robot to achieve. In this problem, you will move the robot to a series of positions, record the joint angles necessary to achieve each position, and then command the robot to repeat these actions.

To record joint angles, run the `teach_waypoints.m` function. This function will repeatedly query feedback from the robot, and when 'enter' is pressed it will save the current joint angles into a matrix; one column for each waypoint. Press 'x' and then 'enter' when you are done training waypoints.

 (1) [15 points] Running through waypoints

You will need to fill in the `repeat_waypoints.m` function. This function will take in a robot object, the matrix of joint angles generated by `teach_waypoints.m`, and a time in seconds used to complete each individual motion (waypoint to waypoint).

In this function, you will use the robot's current location and compute a linear joint trajectory that moves from this location to the first joint angle position in the matrix. You will then create linear joint-space trajectories between each consecutive waypoint.

Finally, you will play through this entire series (starting at the robot's initial posiiton until reaching the final waypoint). control the robot at approximately 100Hz, and therefore will need to create trajectories with the correct number of points in order to complete each motion in the given time.

At the end of the function, two plots will be created for you. First, the workspace path of the end effector, based on your kinematics, and second the joint positions, velocities, and torques.

To get the points for this question, you will need to **demo** the code for the TAs; more information about scheduling this demo will be sent out by the TAs.

In this demo, you will need to teach the robot a series of waypoints (e.g., tracing the star on the table), and then repeat these waypoints with a waypoint-to-waypoint time of 2 seconds. Also, be prepared to discuss the plots that are generated, especially the joint torque plot (areas of high magnitude and how these affect position error).

2) "Drawing" Shapes

A second way to control the behavior of real robots is to be given a set of desired *workspace* positions, and determine how to move the robot to these points. This will be especially relevant for your final project - when placing a series of Jenga blocks into a tower, imagine

you had the block locations encoded as joint angle waypoints, and then wanted to shift the tower over by an inch, or add 5 more layers? This would require training many more waypoints! If you use IK to generate the joint angles for a pattern of x,y,z locations, then you just need to add an offset to the desired location, or increment your loop counter used when computing these angles.

(1) [15 points] Running through waypoints

You will need to fill in the `play_through_workspace_waypoints.m` function. This will take in a robot object, a matrix of workspace positions (each column is an $\begin{bmatrix} x \\ y \end{bmatrix}$ location), and a time in seconds used to complete each individual motion (location to location).

In this function, you will use the robot's current location and compute a straight-line workspace trajectory that moves from this location to the first workspace position in the matrix. You will then calculate straight-line workspace trajectories that move between each workspace position in turn, until reaching the last one.

Finally, you will play through this entire series (starting at the robot's initial posiiton until reaching the final workspace position). You will control the robot at approximately 100Hz, and therefore will need to create trajectories with the correct number of points in order to complete each motion in the desired time.

At the end of the function, two plots will be created for you. First, the workspace path of the end effector, based on your kinematics, and second the joint positions, velocities, and torques.

To get the points for this question, you will need to **demo** the code for the TAs; more information about scheduling this demo will be sent out by the TAs.

In this demo, you will be given a series of workspace positions, and use a position-to-position time of 2 seconds. You will use your function to play through these positions. Be prepared to discuss the plots that are generated, especially the joint torque plot (areas of high magnitude and how these affect position error).

3) Submission

To submit, run `create_submission.m`. It will first check that your submitted files run without error, and perform a small sanity check. Note, this is not going to grade your submission! The function will create a file called `handin-4.tar.gz`. Upload it to Autolab to complete the submission.

7 Submission Checklist

☐ Create a PDF of your answers to the written questions with the name `writeup.pdf`.

☐ Make sure you have `writeup.pdf` in the same folder as the `create_submission.m` script.

☐ Run `create_submission.m` in Matlab.

☐ Upload `handin-4.tar.gz` to Canvas.

☐ After completing the entire assignment, fill out the feedback form[3].

---

[3] https://canvas.cmu.edu/courses/6365/quizzes/14005